

Shared Memory Control Parallelism: OpenMP

Clay Breshears
Intel



Agenda

What is OpenMP?

Parallel regions

Worksharing Constructs

Data environment

Synchronization



What Is OpenMP?

- Portable, shared-memory threading API
 - Fortran, C, and C++
 - Multi-vendor support for both Linux and Windows
 - Standardizes task & loop-level parallelism
 - Supports task-based parallelism
 - Combines task-based and loop-based parallelism
 - Standardizes task-based parallelism
 - Standardizes task-based parallelism
- <http://www.openmp.org>**

Current spec is OpenMP 3.0

318 Pages

<http://www.openmp.org>

Current spec is OpenMP 3.0

318 Pages


(combined C/C++ and Fortran)



Programming Model

Fork-Join Parallelism:

- **Master thread** spawns a **team of threads** as needed
- Parallelism is added incrementally: that is, the sequential program evolves into a parallel program



A Few Syntax Details to Get Started

- Most of the constructs in OpenMP are compiler directives or pragmas
 - For C and C++, the pragmas take the form:

```
#pragma omp construct [clause [clause]...]
```



Parallel Region & Structured Blocks (C/C++)

OpenMP constructs apply to structured blocks

Structured block: a block with one point of entry at the top and one point of exit at the bottom

```
#pragma omp parallel
{
    int id = omp_get_thread_num();
more: res[id] = do_big_job (id);

    if (conv (res[id]) goto more;
}
printf ("All done\n");
```

A structured block

```
if (go_now()) goto more;
#pragma omp parallel
{
    int id = omp_get_thread_num();
more:  res[id] = do_big_job(id);
    if (conv (res[id]) goto done;
    goto more;
}
done: if (!really_done()) goto more;
```

Not a structured block



Worksharing

- **Worksharing** is the general term used in OpenMP to describe distribution of work across threads.
- Three examples of worksharing in OpenMP are:
 - `omp single` construct
 - `omp for` construct
 - `omp task` construct

Automatically divides work
among threads



Single Construct

- Denotes block of code to be executed by only one thread
 - First thread to arrive is chosen
- Implicit barrier at end

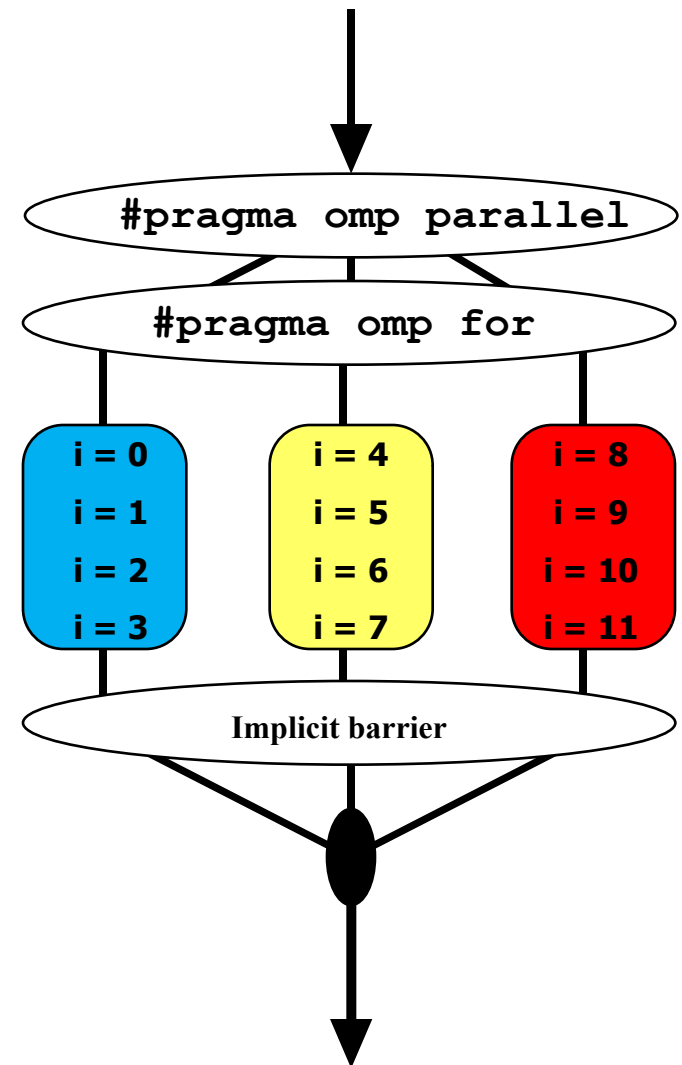
```
#pragma omp parallel
{
    DoManyThings();
    #pragma omp single
    {
        ExchangeBoundaries();
    } // threads wait here for single
    DoManyMoreThings();
}
```



omp for Construct

```
// assume N = 12
#pragma omp parallel
#pragma omp for
    for(i = 0; i < N; i++)
        c[i] = a[i] + b[i];
```

- Threads are assigned an independent set of iterations
- Threads must wait at the end of work-sharing construct



Combining constructs

- These two code segments are equivalent

```
#pragma omp parallel
{
    #pragma omp for
    for (i=0; i< MAX; i++) {
        res[i] = huge();
    }
}
```

```
#pragma omp parallel for
    for (i=0; i< MAX; i++) {
        res[i] = huge();
    }
```



The schedule clause

The **schedule clause** affects how loop iterations are mapped onto threads

`schedule(static [, chunk])`

- Blocks of iterations of size “chunk” to threads
- Round robin distribution
- Low overhead, may cause load imbalance

`schedule(dynamic [, chunk])`

- Threads grab “chunk” iterations
- When done with iterations, thread requests next set
- Higher threading overhead, can reduce load imbalance

`schedule(guided [, chunk])`

- Dynamic schedule starting with large block
- Size of the blocks shrink; no smaller than “chunk”



Schedule Clause Example

```
#pragma omp parallel for schedule (static, 8)
  for( int i = start; i <= end; i += 2 )
  {
    if ( TestForPrime(i) )  gPrimesFound++;
  }
```

Iterations are divided into chunks of 8

- If start = 3, then first chunk is $i=\{3,5,7,9,11,13,15,17\}$



Data Scoping – What's shared

- OpenMP uses a shared-memory programming model
- **Shared variable** - a *variable* whose name provides access to a the same block of storage for each task region
 - Shared clause can be used to make items explicitly shared
 - Global variables are shared among tasks
 - C/C++: File scope variables, namespace scope variables, static variables, Variables with const-qualified type having no mutable member are shared, Static variables which are declared in a scope inside the construct are shared.



Data Scoping – What's private

- But, not everything is shared...
 - Examples of implicitly determined private variables:
 - Stack (local) variables in functions called from parallel regions are PRIVATE
 - Automatic variables within a statement block are PRIVATE
 - Loop iteration variables are private
 - Implicitly declared private variables within tasks will be treated as firstprivate




A Data Environment Example

```
float A[10];
main ()
{
    integer index[10];
    #pragma omp parallel
    {
        Work (index);
    }
    printf ("%d\n", index[1]);
}
```

A, *index*, and *count* are shared by all threads, but *temp* is local to each thread

```
extern float A[10];
void Work (int *index)
{
    float temp[10];
    static integer count;
    <...>
}
```



The Private Clause

- Reproduces the variable for each task
 - Variables are un-initialized; C++ object is default constructed
 - Any value external to the parallel region is undefined

```
void* work(float* c, int N) {  
    float x, y; int i;  
    #pragma omp parallel for private(x,y)  
        for(i=0; i<N; i++) {  
            x = a[i]; y = b[i];  
            c[i] = x + y;  
        }  
}
```




New Addition to OpenMP

- **Tasks** – Main change for OpenMP 3.0
- Allows parallelization of irregular problems
 - unbounded loops
 - recursive algorithms
 - producer/consumer



What are tasks?

- Tasks are independent units of work
- Threads are assigned to perform the work of each task
 - Tasks may be deferred
- Tasks may be executed immediately
- The runtime system decides which of the above
 - Tasks are composed of:
 - **code** to execute
 - **data** environment
 - **internal control variables** (ICV)



Task Construct – Explicit Task View

- A team of threads is created at the **omp parallel** construct
- A single thread, T0, is chosen to execute the **while** loop
- T0 operates the **while** loop, creates tasks, and fetches next pointers
- Each time T0 crosses the **omp task** construct it generates a new task
- Each task runs in its own thread
- All tasks complete at the barrier at the end of the parallel region's single construct

```
#pragma omp parallel
{
    #pragma omp single
    { // block 1
        node * p = head;
        while (p) { //block 2
            #pragma omp task
            process(p);
            p = p->next; //block 3
        }
    } // tasks done
}
```




Why are tasks useful?

Have potential to parallelize irregular patterns and recursive function calls

```
#pragma omp parallel
{
  #pragma omp single
  { // block 1
    node * p = head;
    while (p) { //block 2
      #pragma omp task
      process(p);
      p = p->next; //block 3
    }
  }
}
```

Single
Threaded







Thr1

Thr2

Thr3

Thr4



Time

Time Saved



When are tasks guaranteed to be complete?

Tasks are guaranteed to be complete:

- At thread or task barriers
- At the directive: `#pragma omp barrier`
- At the directive: `#pragma omp taskwait`



Example: Dot Product

```
float dot_prod(float* a, float* b, int N)
{
    float sum = 0.0;
    #pragma omp parallel for
    for(int i=0; i<N; i++) {
        sum += a[i] * b[i];
    }
    return sum;
}
```

What is Wrong?





Race Condition

- A *race condition* is nondeterministic behavior caused by the times at which two or more threads access a shared variable
- For example, suppose both Thread A and Thread B are executing the statement

```
area += 4.0 / (1.0 + x*x) ;
```



Two Timings



Protect Shared Data

- Must protect access to shared, modifiable data

```
float dot_prod(float* a, float* b, int N)
{
    float sum = 0.0;
    #pragma omp parallel for
        for(int i=0; i<N; i++) {
    #pragma omp critical
        sum += a[i] * b[i];
    }
    return sum;
}
```




OpenMP Critical Construct

```
#pragma omp critical [(lock_name)]
```

- Defines a critical region on a structured block

Threads wait their turn –at a time, only one calls `consum()` thereby protecting `RES` from race conditions

Naming the critical construct `RES_lock` is optional



```
float RES;  
#pragma omp parallel  
{ float B;  
#pragma omp for  
  for(int i=0; i<niters; i++){  
    B = big_job(i);  
#pragma omp critical (RES_lock)  
    consum (B, RES);  
  }  
}
```

Good Practice – Name all critical sections



OpenMP* Reduction Clause

reduction (*op* : *list*)

- The variables in “*list*” must be shared in the enclosing parallel region
- Inside parallel or work-sharing construct:
 - A PRIVATE copy of each list variable is created and initialized depending on the “*op*”
 - These copies are updated locally by threads
 - At end of construct, local copies are combined through “*op*” into a single value and combined with the value in the original SHARED variable



Reduction Example

```
#pragma omp parallel for reduction(+:sum)
for(i=0; i<N; i++) {
    sum += a[i] * b[i];
}
```

- Local copy of *sum* for each thread
- All local copies of *sum* added together and stored in “global” variable



C/C++ Reduction Operations


- A range of associative operands can be used with reduction
- Initial values are the ones that make sense mathematically

Operand	Initial Value
+	0
*	1
-	0
^	0

Operand	Initial Value
&	~ 0
	0
&&	1
	0



Numerical Integration Example



```
static long num_steps=100000;  
double step, pi;
```

```
void main()  
{  int i;  
    double x, sum = 0.0;  
  
    step = 1.0/(double) num_steps;  
    for (i=0; i< num_steps; i++){  
        x = (i+0.5)*step;  
        sum = sum + 4.0/(1.0 + x*x);  
    }  
    pi = step * sum;  
    printf("Pi = %f\n",pi);  
}
```



Computing Pi

```
static long num_steps=100000;
double step, pi;

void main()
{  int i;
   double x, sum = 0.0;
   step = 1.0/(double) num_steps;
   #pragma omp parallel for private(x) reduction(+:sum)
   for (i=0; i< num_steps; i++){
       x = (i+0.5)*step;
       sum = sum + 4.0/(1.0 + x*x);
   }
   pi = step * sum;
   printf("Pi = %f\n",pi);
}
```

can be

need

should



Atomic Construct

- Special case of a critical section
- Applies only to simple update of memory location

```
#pragma omp parallel for shared(x, y, index, n)
  for (i = 0; i < n; i++) {
    #pragma omp atomic
      x[index[i]] += work1(i);
      y[i] += work2(i);
  }
```



Master Construct

- Denotes block of code to be executed only by the master thread
- No implicit barrier at end

```
#pragma omp parallel
{
    DoManyThings();
    #pragma omp master
    {
        // if not master skip to next stmt
        ExchangeBoundaries();
    }
    DoManyMoreThings();
}
```



Implicit Barriers

- Several OpenMP* constructs have implicit barriers
 - Parallel – necessary barrier – cannot be removed
 - for
 - single
- Unnecessary barriers hurt performance and can be removed with the nowait clause
 - The nowait clause is applicable to:
 - For clause
 - Single clause



Nowait Clause

```
#pragma omp for nowait
for(...)
{...};
```

```
#pragma single nowait
{ [...] }
```

- Use when threads unnecessarily wait between independent computations

```
#pragma omp for schedule(dynamic,1) nowait
for(int i=0; i<n; i++)
    a[i] = bigFunc1(i);

#pragma omp for schedule(dynamic,1)
for(int j=0; j<m; j++)
    b[j] = bigFunc2(j);
```



Barrier Construct

- Explicit barrier synchronization
- Each thread waits until all threads arrive

```
#pragma omp parallel shared (A, B, C)
{
    DoSomeWork (A,B) ;
    printf("Processed A into B\n") ;
#pragma omp barrier
    DoSomeWork (B,C) ;
    printf("Processed B into C\n") ;
}
```

